

TCS 802 Advanced Computer Architecture

Pipelined Processing & Superscalar Techniques

Linear & Nonlinear pipelines
Instruction Pipelines & Arithmetic Operations
Superscalar and super-pipeline

Principles of Pipelining

- A pipeline may be compared directly with an assembly line in a manufacturing plant. Thus
 - Input task or process is subdivided into a sequence of subtasks;
 - Each subtask is executed by a specialized hardware stage;
 - Many such hardware stages operate concurrently;
 - When successive tasks are streamed into the pipeline they are executed in an overlapped fashion at the subtask level.
 - The creation of the correct sequence of subtasks is crucial to the performance of the pipeline.
 - Slowest subtask is the bottleneck in the pipeline.

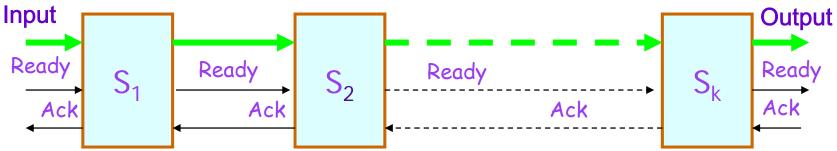
Linear Pipeline



- Processing Stages are linearly connected
- Perform fixed function
- Synchronous Pipeline
 - Clocked latches between Stage i and Stage i+1
 - Equal delays in all stages
- Asynchronous Pipeline (Handshaking)

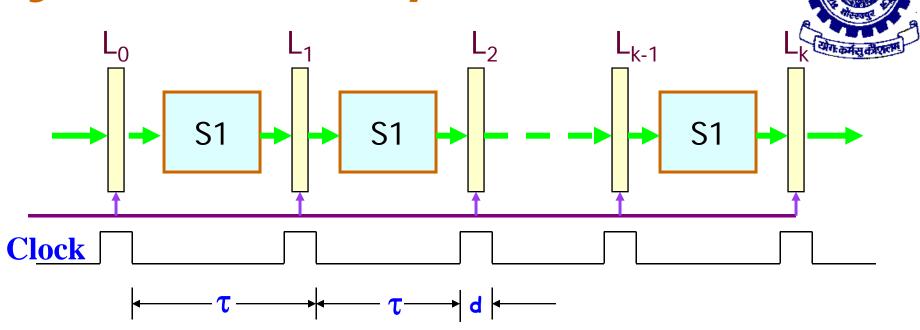
Asynchronous Pipeline



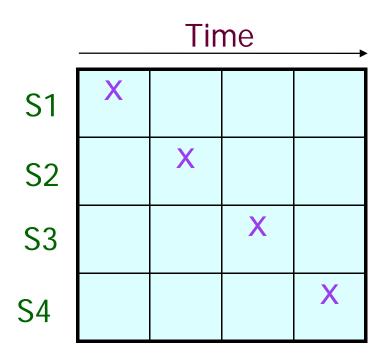


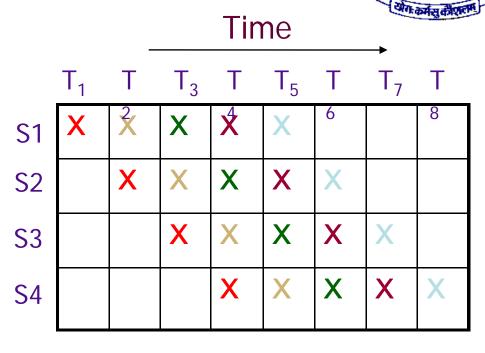
- Asynchronous pipelines are useful in designing communication channels in massage passing multicomputer
- Asynchronous pipelines may have variable throughput rate. Different amount of delay may be experienced in different stages.

Synchronous Pipeline



Reservation Table





Reservation Table of a four stage linear pipeline

5 tasks on 4 stages

Clock period and frequency



Define

 τ_i to be time delay due to logic circuitry in stage S_i , d to be time delay of each interface latch.

Then

• The clock period, τ , of a linear pipeline is given by

$$\tau = \max \{\tau_i\} + d$$
$$= \tau_M + d$$

• The frequency, $f = 1 / \tau$ = 1 / [τ_M + d] (i.e., reciprocal of clock period)

Speedup, Efficiency & Throughput



Speedup

Ideally, a linear pipeline of k stages can process n task in k + (n-1) clock so total time required is

$$T_K = \{ k + (n-1) \} \tau$$

Time required for nonpipelined processing

$$T_1 = n k \tau$$

Speedup factor
$$S_K = T_1 / T_k$$

$$S_K = n k / k + (n-1)$$

Speedup, Efficiency & Throughput



- Pipeline efficiency, η = S_k / k
 = n / k + (n-1)
 So efficiency = 1 when n→∞ and
 lower bound of η is 1 / k when n=1
- Throughput i.e. Number of task performed per unit time

$$H_k = n / \{ k + (n-1) \} \tau = n f / k + (n-1)$$

Non Linear Pipelines



- Variable functions
- Feed-Forward
- Feedback

Linear Instruction Pipelines



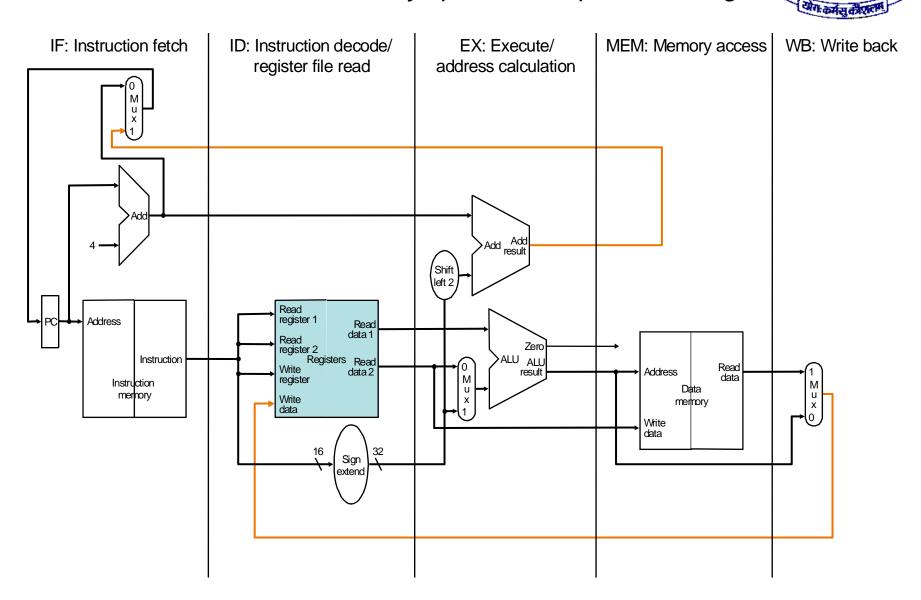
- Assume the following instruction execution phases:
 - Fetch (F)
 - Decode (D)
 - Operand Fetch (O)
 - Execute (E)
 - Write results (W)

Execution

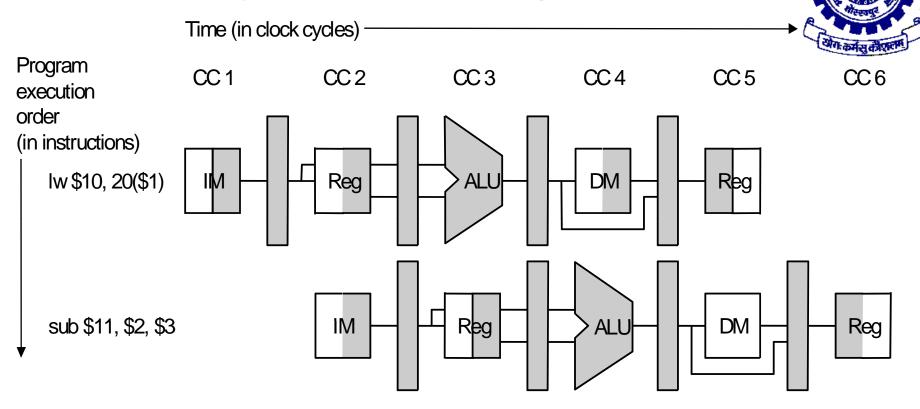
F	1	12	I ₃					
D		I ₁	12	13				
O			11	12	I ₃			
E				I ₁		I ₃		
W					1	12	I ₃	

Basic idea: start from single cycle impl.

What do we need to add to actually split the datapath into stages?



Graphically Representing Pipelines



- Can help with answering questions like:
 - how many cycles does it take to execute this code?
 - what is the ALU doing during cycle 4?
 - use this representation to help understand datapaths

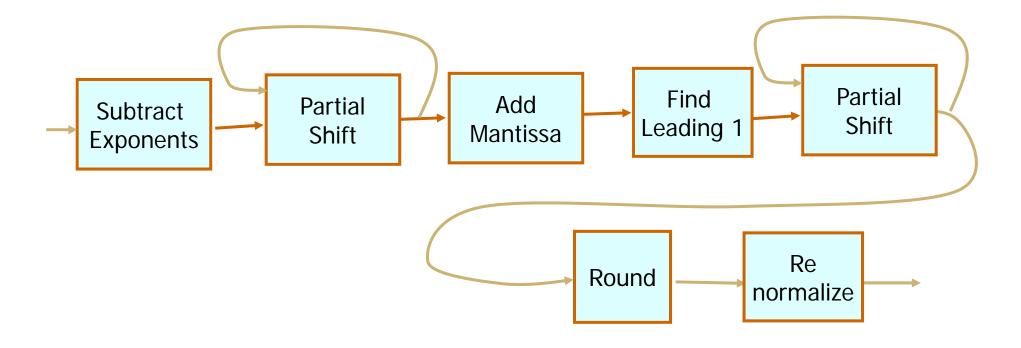
Floating Point Multiplication



- Inputs (Mantissa₁, Exponenet₁), (Mantissa₂, Exponent₂)
- ◆ Add the two exponents → Exponent-out
- Multiple the 2 mantissas
- Normalize mantissa and adjust exponent
- Round the product mantissa to a single length mantissa. You may adjust the exponent

Linear Pipeline for floating-point Addition

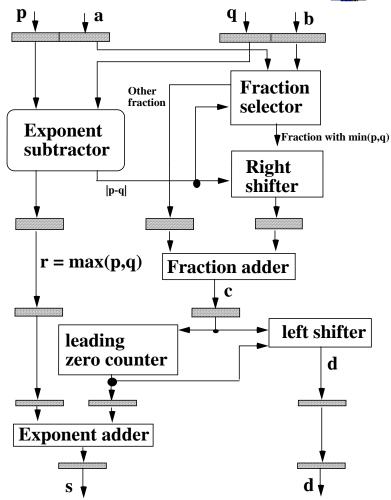




Example

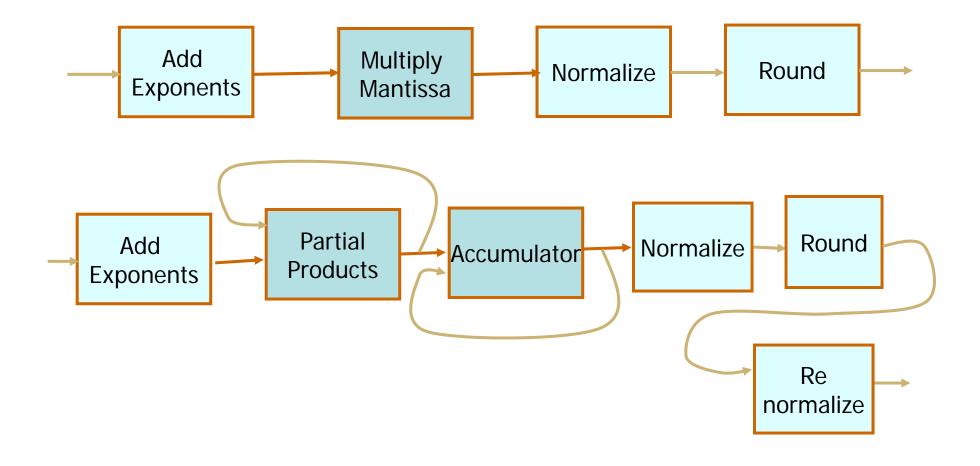
- Linear pipeline with four functional stages. Inputs are two normalised floating-point numbers
 a*2^p and b*2^q
- Output is a normalised floatingpoint number **d* 2^s** which is the sum of the two inputs.
- The hardware units other than the latches can all be implemented using combinational logic.
- If time delay of interface latches is 10ns and if the time delays of the four stages are 60, 50, 90 and 80ns, respectively, then cycle time of pipeline can be chosen to be 100ns.





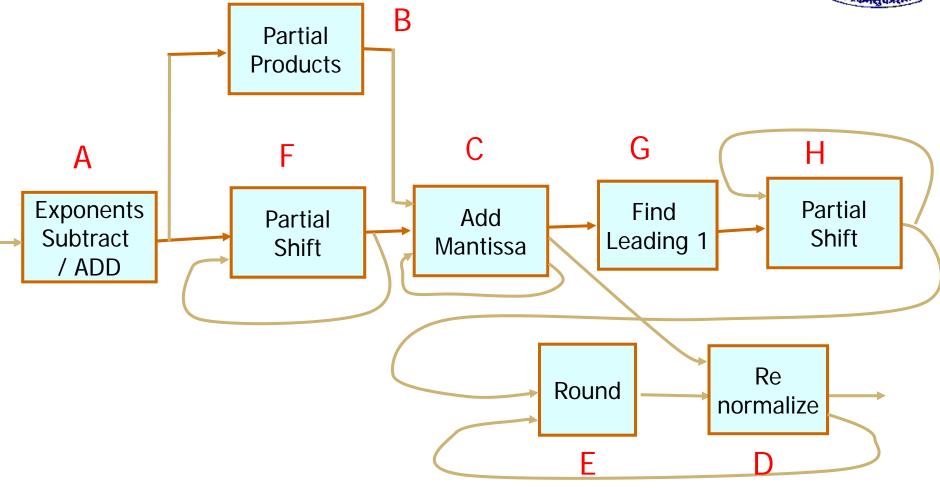
Linear Pipeline for floating-point multiplication





Combined Adder and Multiplier





Reservation Table for Multiply



	1	2	3	4	5	6	7
Α	X						
В		Х	X				
С			Х	Х			
D					Х		X
E						X	
F							
G							
Н							

Reservation Table for Addition

	1	2	3	4	5	6	7	8	9
Α	Υ								
В									
С				Y					
D									Y
Е								Y	
F		Y	Y						
G					Y				
Н						Y	Y		

Nonlinear Pipeline Design



- Latency
 - The number of clock cycles between two initiations of a pipeline
- Collision
 - Resource Conflict
- Forbidden Latencies
 - Latencies that cause collisions

Nonlinear Pipeline Design



- Latency Sequence
 - A sequence of permissible latencies between successive task initiations
- Latency Cycle
 - A sequence that repeats the same subsequence
- Collision vector
 - **□** C = (Cm, Cm-1, ..., C2, C1), m <= n-1
 - n = number of column in reservation table
 - Ci = 1 if latency i causes collision, 0 otherwise

Mul – Mul Collision (lunch after 1 cycle)

	1	2	3	4	5	6	7
Α	X	Z					
В		X	ΧZ	Z			
C			X	ΧZ	Z		
D					X	Z	X
Е						X	Z
F							
G							
Н							

Mul – Mul Collision (lunch after 2 cycle)



	1	2	3	4	5	6	7
Α	X		Z				
В		X	X	Z	Z		
С			X	X	Z	Z	
D					X		XZ
Е						X	
F							
G							
Н							

Mul – Mul Collision (lunch after 3 cycle)



	1	2	3	4	5	6	7
Α	X			Z			
В		X	X		Z	Z	
С			X	X		Z	Z
D					X		X
Е						X	
F							
G							
Н							

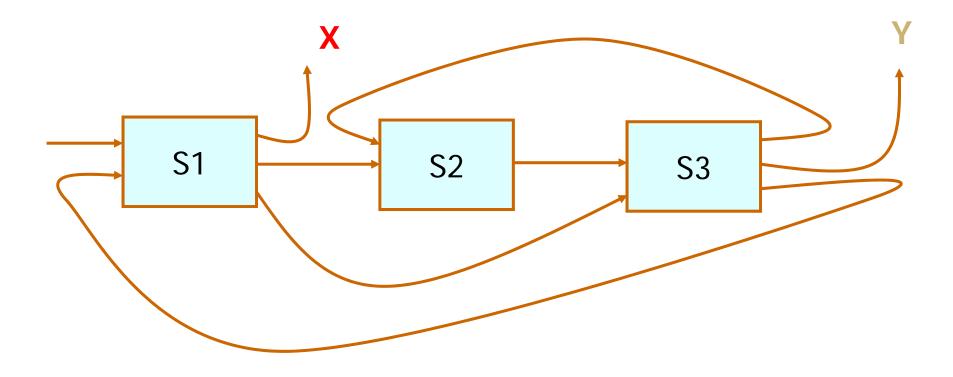
Collision Vector for Multiply after Multiply



- Forbidden Latencies: 1, 2
- Collision vector
 0 0 0 0 1 1 → 11
- Maximum forbidden latency = 2 → m = 2

Design Example





Reservation Tables for X & Y



S1	X					X		X
S2		X		X				
S 3			X		X		X	

 S1
 Y

 S2
 Y

 S3
 Y

Forbidden Latencies

and the afterest

X after X

2

S 1	X1	X2		X1	X2 X1
------------	----	----	--	----	-------

S2 X1 X2 X1 X2

S3 X1 X2 X1 X2 X1

5

S1	X1			X2 X1	X1
31					

S2 X1 X1 X2

S 3		X1	X1	X1	X2

Forbidden Latencies



X after X

	4				_			
S1	X1				X2	X1		X1
S2		X1		X1		X2		X2
S 3			X1		X1		X2 X1	
	7							
_	V1				Y	1	V	2 V1

S3

S2

	X1		X1			
		X1		X1	X1	

Collision Vector



- Forbidden Latencies: 2, 4, 5, 7
- Collision Vector =1 0 1 1 0 1 0

Y after Y



S 1	Y		Y		Y	
S 2			Y		Υ	
S 3		Υ		YY		YY
				1		
S 1	Y				YY	
S 2			Υ			

S3

Collision Vector



- Forbidden Latencies: 2, 4
- Collision Vector =

1010

Exercise – Find the collision vector

	1	2	3	4	5	6	7
A	X		Χ	X			
В		X				X	
С					X		X
D				X			

Latency Analysis



- The number of cycles between two initiations is the *latency* between them
- ♠ A latency of k → two initiations are separated by k cycles
- Collision > resource conflict between two initiations
- ♠ Latencies that cause collision → forbidden latencies

Latency Analysis



- ♣ Latency Sequence → a sequence of permissible latencies between successive initiations
- ♣ Latency Cycle → a latency sequence that repeats the same subsequence (cycle) indefinitely
- ♣ Latency Sequence → 1, 8
- Latencies Cycle → (1,8) → 1, 8, 1, 8, 1, 8 ...

Latency Analysis

- ♣ Average Latency (of a latency cycle) → sum of all latencies / number of latencies along the cycle
- Constant Cycle → One latency value

Latency Cycle (1,8)



 	 						•		ļ										-		
1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20	21	
X1	X2				X1	X2	X1	X2	X3	X4				X3	X4	X3	X4	X5	X6		
	X1	X2	X1	X2						X3	X4	X3	X4						X5	X6	
		X1	X2	X1	X2	X1	X2				X3	X4	X3	X4	X3	X4				X5	

Average Latency = (1+8)/2 = 4.5

Latency Cycle (6)



 				,	ļ						 							1			
1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20	21	
X1					X1	X2	X1				X2	X3	X2				X 3	X4	X3		
	X1		X1				X2		X2				X3		X3				X4		
		X1		X1		X1		X2		X2		X2		X3		X3		X3		X4	

Average Latency = 6

Find the collision vector



	1	2	3	4	5	6	7
Α	X		X	X			
В		X				X	
С					X		X
D				X			

Collision Vector

- Forbidden Latencies: 1,2, 4
- Collision Vector =

001011

Reservation Tables for X & Y



S1	X					X		X
S2		X		X				
S 3			X		X		X	

S1	Y				Y	
S2			Y			
S 3		Y		Y		Υ

Collision Vector

जिया के स्थाप के स्थ

- Forbidden Latencies: 2, 4, 5, 7
- Collision Vector =

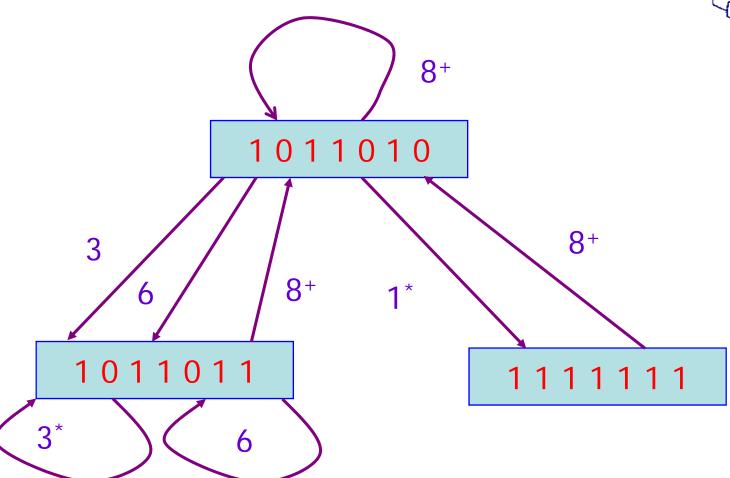
1011010

- The next state is obtained by bitwise ORing the initial collision vector with the shifted register
 - C.V. = 1 0 1 1 0 1 0 (first state)

```
0 1 0 1 1 0 1 C.V. 1-bit right shifted
1 0 1 1 0 1 0 initial C.V.
----- OR
1 1 1 1 1 1 1
```

State Diagram for X after X





Cycles



- Simple cycles → each state appears only once (3), (6), (8), (1, 8), (3, 8), and (6,8)
- Greedy Cycles → simple cycles whose edges are all made with minimum latencies from their respective starting states

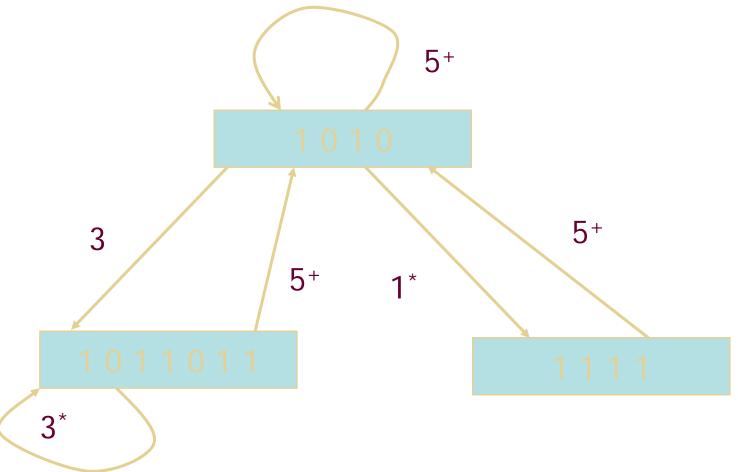
(1,8), $(3) \rightarrow$ one of them is MAL

MAL

- Minimum Average latency
- At least one of the greedy cycles will lead to the MAL
- Consider state diagram for Y, MAL is 3 (See diagram)

State Diagram for Y





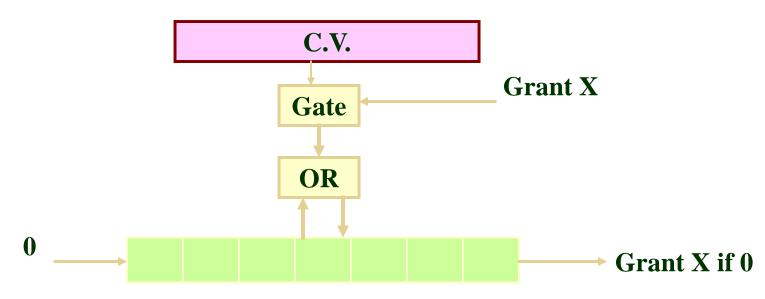
Bounds on the MAL

- Tall to both and the state of t
- MAL is lower bounded by the maximum number of checkmarks in any row of the reservation table. (Shar, 1972)
- MAL is lower than or equal to the average latency of any greedy cycle in the state diagram. (Shar, 1972)
- The average latency of any greedy cycle is upperbounded by the number of 1's in the initial collision vector plus 1. This is also an upper bund on the MAL. (Shar, 1972)

Single Function Controller

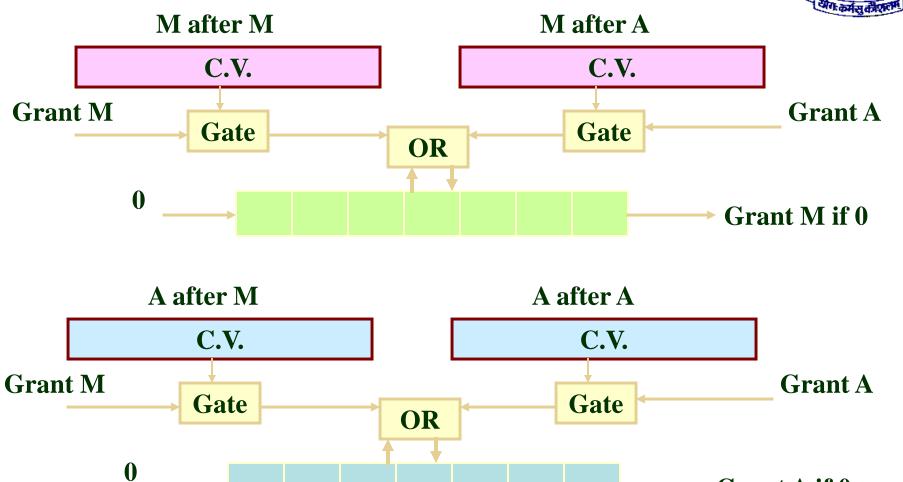


X after X



Controller for a dual-function pipeline





Grant A if 0

State Diagram



- It specifies the permissible state transitions among successive initiations
- Collision vector corresponds to the initial state at time
 t = 1 (initial collision vector)
- The next state comes at time t + p, where p is a permissible latency in the range 1 <= p < m</p>

Right Shift Register



The next state can be obtained with the help of an m-bit shift register



Each 1-bit shift corresponds to increase in the latency by 1

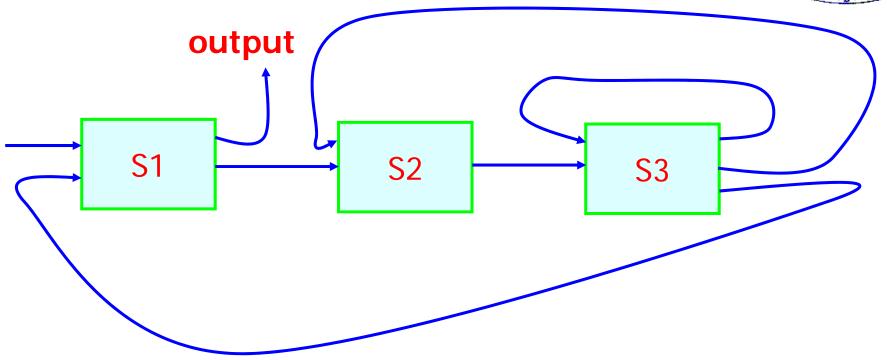
Delay Insertion



- The purpose is to modify the reservation table, yielding a new collision vector
- This may lead to a modified state diagram, which may produce greedy cycles meeting the lower bound on MAL

Example





Reservation Tables

	1	2	3	4	5
S1	X				X
S2		X		X	
S3			X	X	



Forbidden Latencies: 1, 2, 4

5+

 $C.V. \rightarrow 1011$

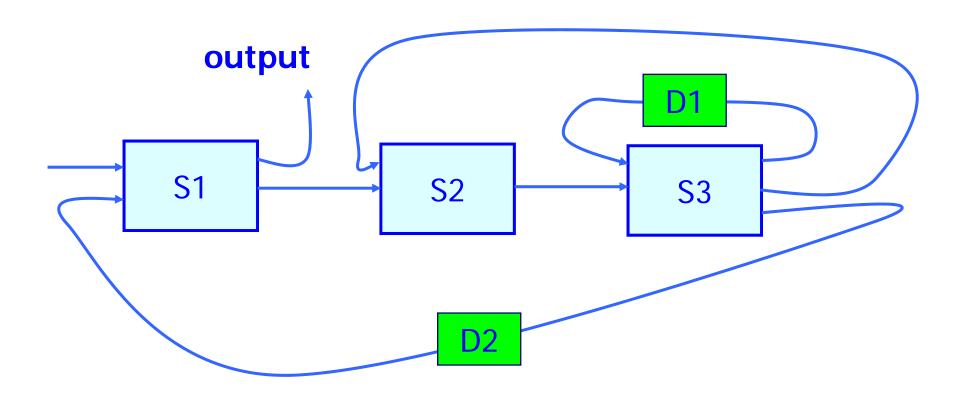
3*

MAL = 3

State Diagram

Example (Cont.)





Example (Cont.)



	1	2	3	4	5	6	7
S1	X						X
S2		X		X			
S3			Х		Х		
D1				X			
D2						X	

Forbidden: 2, 6

 $C.V. \rightarrow 100010$



Hazards



Hazards: problems due to pipelining

Hazard types:

- Structural
 - same resource is needed multiple times in the same cycle
- Data
 - data dependencies limit pipelining
- Control
 - next executed instruction may not be the next specified instruction

Structural hazards



Examples:

- Two accesses to a single ported memory
- Two operations need the same function unit at the same time
- Two operations need the same function unit in successive cycles, but the unit is not pipelined

Solutions:

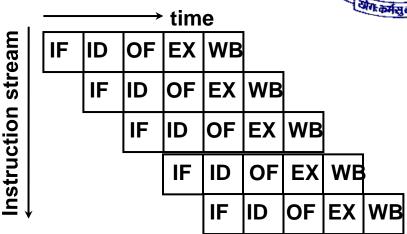
- stalling
- add more hardware

Structural hazards

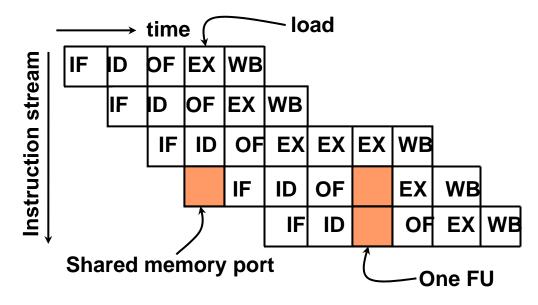
Table and a different

Simple pipelining diagram (but not MIPS!):

- IF: instruction fetch
- ID: instruction decode
- OF: operand fetch
- EX: execute stage(s)
- WB: write back



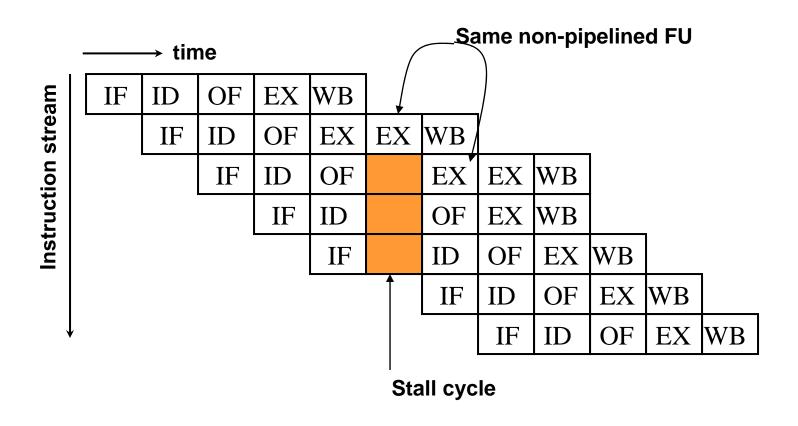
Pipeline stalls due to lack of resources:



Structural hazards



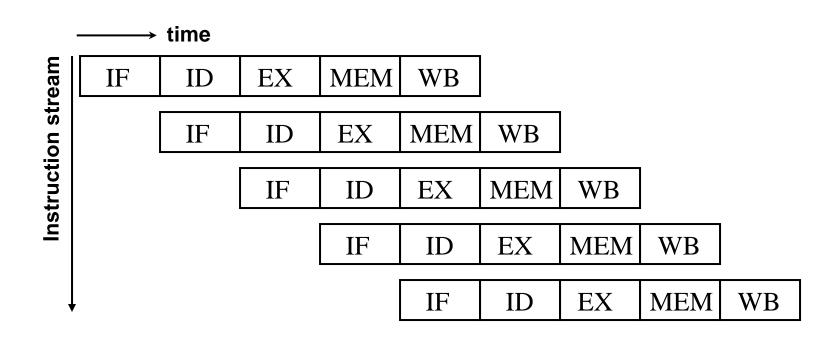
Non-pipelined units



Structural hazards on MIPS



Q: Do we have structural hazards on our simple MIPS pipeline?



Data hazards



- Data dependencies:
 - RaW (read-after-write)
 - WaW (write-after-write)
 - WaR (write-after-read)
- Hardware solution:
 - Forwarding / Bypassing
 - Detection logic
 - Stalling
- Software solution: Scheduling

Data dependences

Three types: RaW, WaR and WaW

```
add r1, r2, 5 ; r1 := r2+5 sub r4, r1, r3 ; RaW of r1

add r1, r2, 5 sub r2, r4, 1 ; WaR of r2

add r1, r2, 5 sub r1, r1, 1 ; WaW of r1

st r1, 5(r2) ; M[r2+5] := r1 ld r5, 0(r4) ; RaW if 5+r2 = 0+r4
```

WaW and WaR do not occur in simple pipelines, but they limit scheduling freedom!

Problems for your compiler and Pentium!

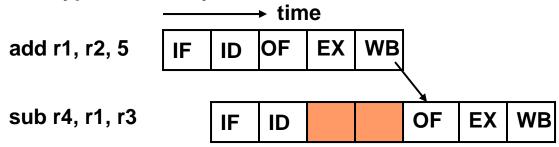
 \Rightarrow use register renaming to solve this!

RaW dependence

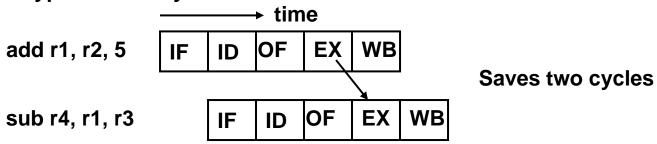


```
add r1, r2, 5 ;r1:= r2+5
sub r4, r1, r3 ;RaW of r1
```

Without bypass circuitry



With bypass circuitry

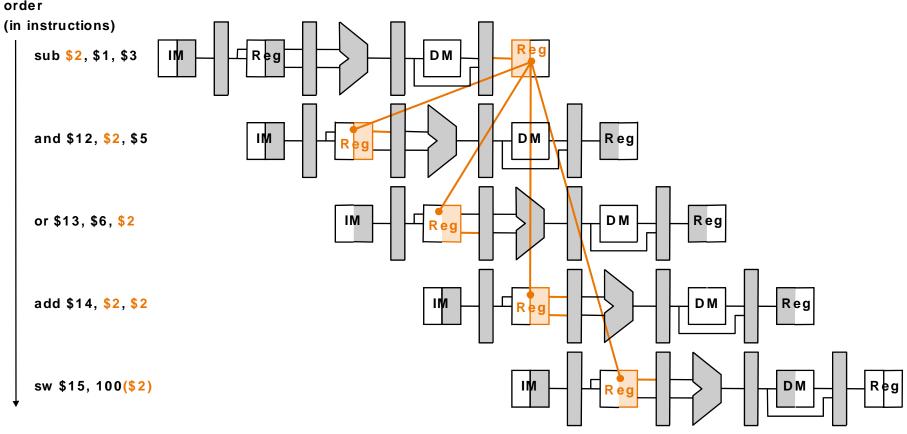


RaW on MIPS pipeline

Time (in clock cycles)

CC 5 CC 6 CC8 CC 9 CC 1 CC 2 CC3 CC 4 CC 7 Value of register \$2: 10 10 10/- 20 -20- 20 10 10 - 20 - 20

Program execution order



Forwarding

Use temporary results, don't wait for them to be written

Χ

register file forwarding to handle read/write to same register

Χ

Value of MEM/WB: X

■ ALU for								योगः कर्मसु कीएलम		
	CC 1	CC 2	CC 3	CC 4	CC 5	CC 6	CC 7	CC 8	CC 9	
Value of register \$2: 10 10			10	10	10/- 20	- 20	- 20	- 20	- 20	
Value of EX/N	ИЕМ: X	X	X	- 20	X	X	X	X	X	

- 20

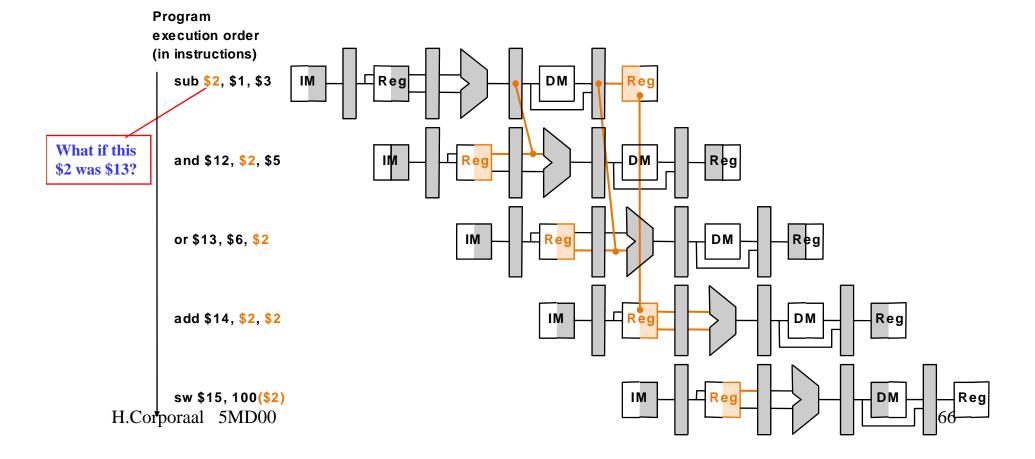
Χ

Χ

Χ

Χ

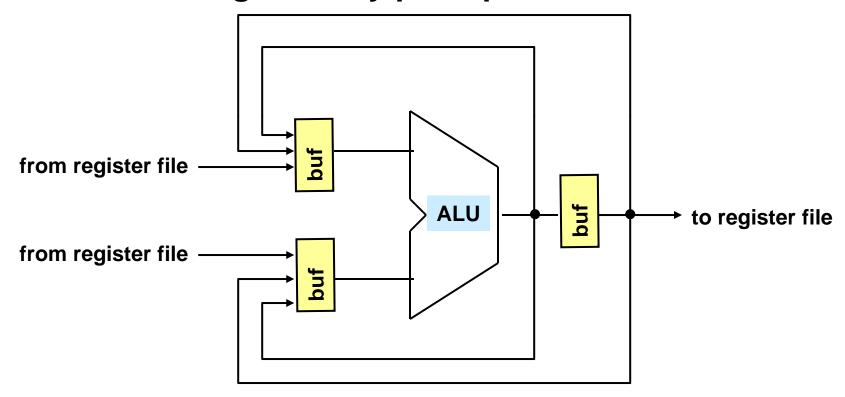
X



Forwarding hardware

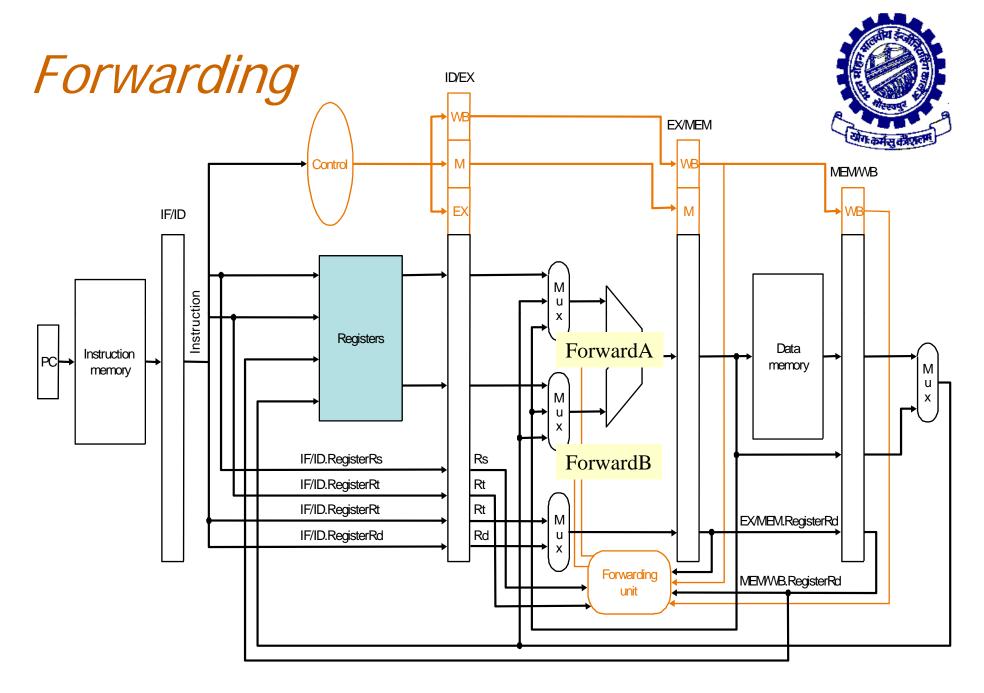


ALU forwarding circuitry principle:



Note: there are two options

- buf ALU bypass mux buf
- buf bypass mux ALU buf



H.Corporaal 5MD00

Forwarding check



- Check for matching register-ids:
- For each source-id of operation in the EX-stage check if there is a matching pending dest-id

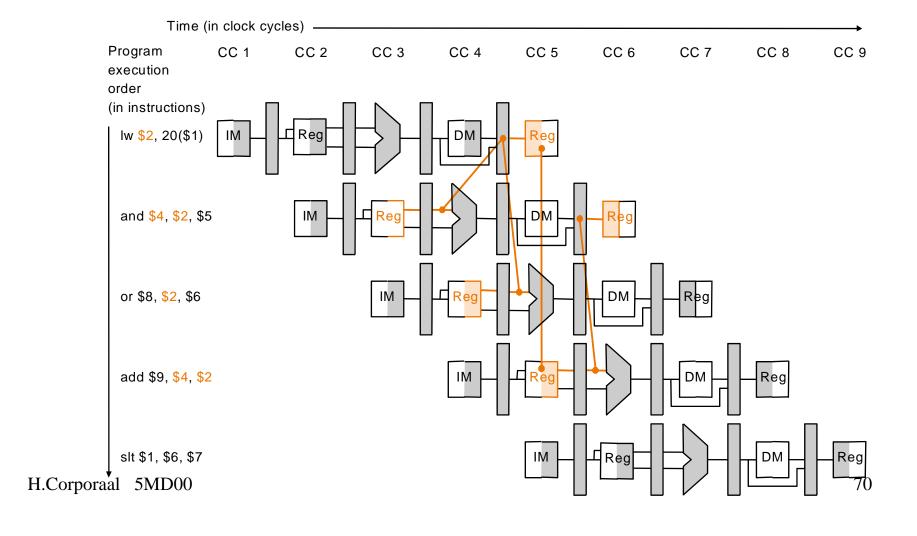
Example:

```
if (EX/MEM.RegWrite) \
   (EX/MEM.RegisterRd \neq 0) \\
   (EX/MEM.RegisterRd = ID/EX.RegisterRs)
then ForwardA = 10
```

Q. How many comparators do we need?

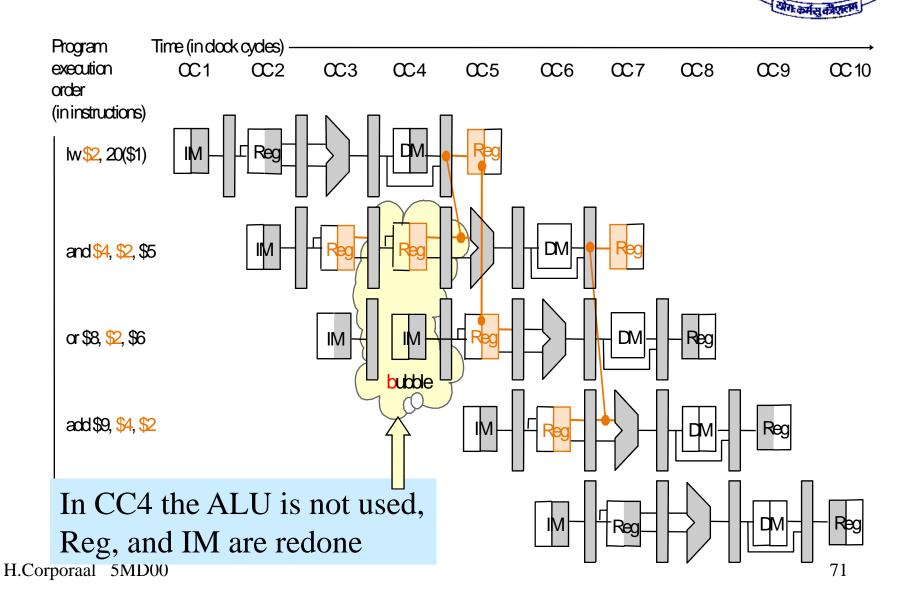
Can't always forward

- Load word can still cause a hazard:
 - an instruction tries to read register **r** following a load to the same
- Need a hazard detection unit to "stall" the load instruction

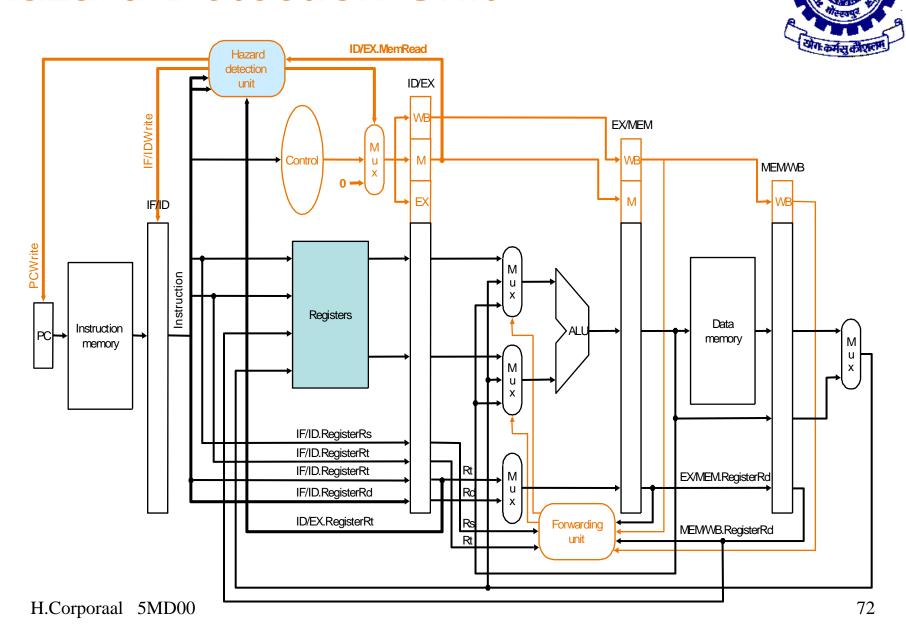


Stalling

We can stall the pipeline by keeping an instruction in the same stage



Hazard Detection Unit



Software only solution?



- Have compiler guarantee that no hazards occur
- Example: where do we insert the "NOPs" ?

Problem: this really slows us down!

sub	\$2,	\$1,	\$3
nop			
nop			
and	\$12,	\$2,	\$5
or	\$13,	\$6,	\$2
add	\$14,	\$2,	\$2
nop			
SW	\$13,	100(\$2)

Control hazards



- Control operations may change the sequential flow of instructions
 - branch
 - jump
 - call (jump and link)
 - return
 - (exception/interrupt and rti / return from interrupt)

H.Corporaal 5MD00

Control hazard: Branch



Branch actions:

- Compute new address
- Determine condition
- Perform the actual branch (if taken): PC := new address

H.Corporaal 5MD00 75

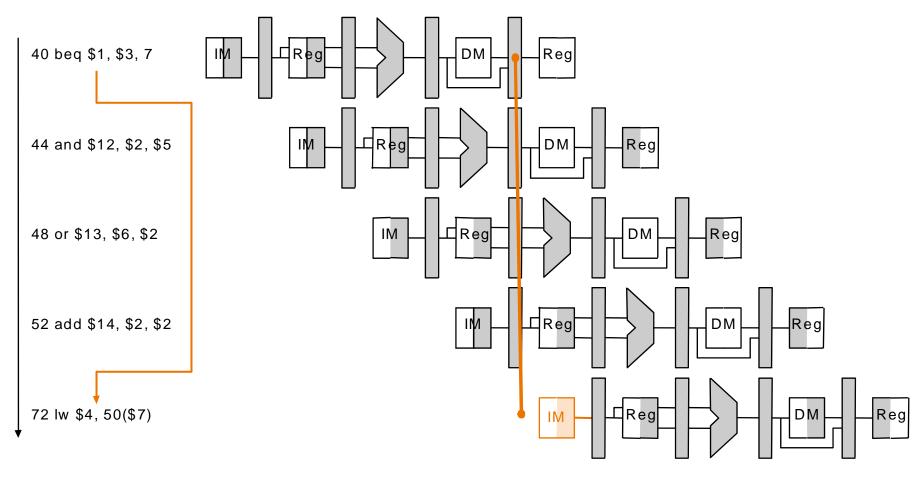
Branch example

Program

execution
order

CC 1 CC 2 CC 3 CC 4 CC 5 CC 6 CC 7 CC 8 CC 9

(in instructions)



Branching



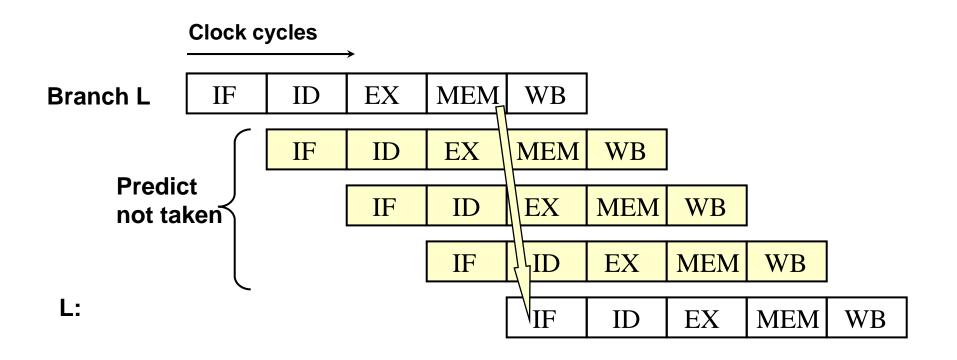
Squash pipeline:

- When we decide to branch, other instructions are in the pipeline!
- We are predicting "branch not taken"
 - need to add hardware for flushing instructions if we are wrong

H.Corporaal 5MD00



Branch with predict not taken

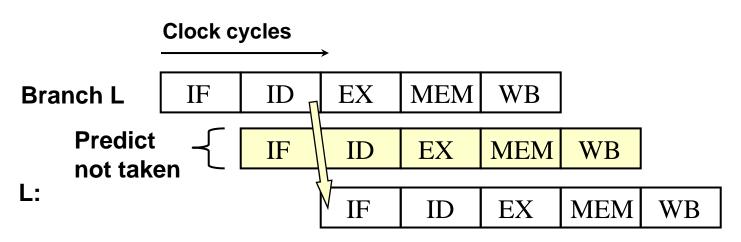


H.Corporaal 5MD00 78

Branch speedup



- Earlier address computation
- Earlier condition calculation
- Put both in the ID pipeline stage
 - adder
 - comparator



Improved branching / flushing IF detection ID/EX EX/MEM MEM/WB IF/ID Registers Data memory memory Forwarding

H.Corporaal 5MD00

80

Exception support



Types of exceptions:

- Overflow
- I/O device request
- Operating system call
- Undefined instruction
- Hardware malfunction
- Page fault
- Precise exception:
 - finish previous instructions (which are still in the pipeline)
 - flush excepting and following instructions, redo them after handling the exception(s)

Exceptions



Changes needed for handling overflow exception of an operation in EX stage (see book for details):

- Extend PC input mux with extra entry with fixed address
- Add EPC register recording the ID/EX stage PC this is the address of the next instruction!
- Cause register recording exception type

E.g., in case of overflow exception insert 3 bubbles; flush the following stages:

- IF/ID stage
- ID/EX stage
- EX/MEM stage

Scheduling, why?



Let's look at the execution time:

$$T_{execution} = N_{cycles} \times T_{cycle}$$

= $N_{instructions} \times CPI \times T_{cycle}$

Scheduling may reduce $T_{execution}$

- Reduce CPI (cycles per instruction)
 - early scheduling of long latency operations
 - avoid pipeline stalls due to structural, data and control hazards
 - allow *N_{issue}* > 1 and therefore *CPI* < 1
- Reduce **N**instructions
 - compact many operations into each instruction (VLIW)

Scheduling data hazards: example 1



Try and avoid RaW stalls (in this case load interlocks)! E.g., reorder these instructions:

```
lw $t0, 0($t1)
lw $t2, 4($t1)
sw $t2, 0($t1)
sw $t0, 4($t1)
```



```
lw $t0, 0($t1)
lw $t2, 4($t1)
sw $t0, 4($t1)
sw $t2, 0($t1)
```

Scheduling data hazards

example 2

Avoiding RaW stalls:

Reordering instructions for following program

(by you or the compiler)



a = b + c

d = e - f



Unscheduled code:

Lw R1,b

Lw R2,c

Add R3,R1,R2 interlock

Sw a,R3

Lw R1,e

Lw R2,f

Sub R4,R1,R2 interlock

Sw d,R4



Scheduled code:

Lw R1,b

Lw R2,c

Lw R5,e extra reg. needed!

Add R3,R1,R2

Lw R2,f

Sw a,R3

Sub R4, R5, R2

Sw d,R4

Scheduling control hazards



$$T_{execution} = N_{instructions} \times CPI \times T_{cycle}$$

$$CPI = CPI_{ideal} + f_{branch} \times P_{branch}$$

$$P_{branch} = N_{delayslots} \times miss_rate$$

- Modern processors tend to have large branch penalty, P_{branch} , due to:
 - many pipeline stages
 - multi-issue
- Note that penalties have larger effect when CPI_{ideal} is low





What can we do about control hazards and CPI penalty?

- Keep penalty P_{branch} low:
 - Early computation of new PC
 - Early determination of condition
 - Visible branch delay slots filled by compiler (MIPS)
- Branch prediction
- Reduce control dependencies (control height reduction) [Schlansker and Kathail, Micro'95]
- Remove branches: if-conversion
 - Conditional instructions: CMOVE, cond skip next
 - Guarding all instructions: TriMedia

Branch delay slot



- Add a branch delay slot:
 - the next instruction after a branch is always executed
 - rely on compiler to "fill" the slot with something useful

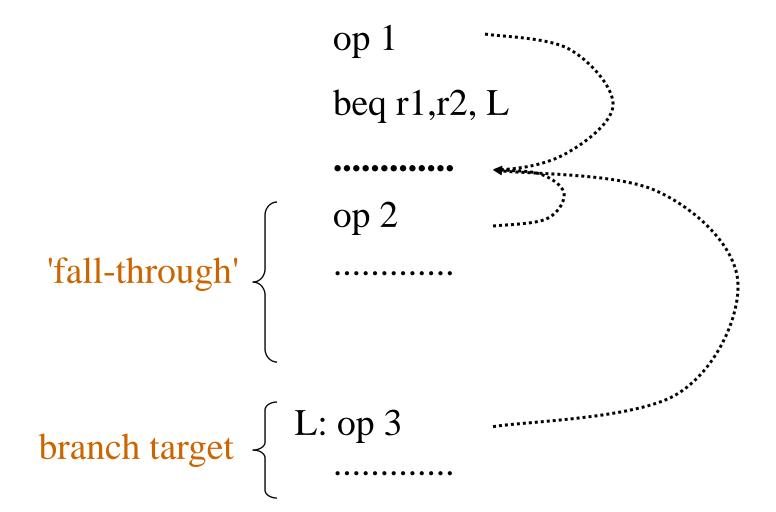
- Is this a good idea?
 - let's look how it works

H.Corporaal 5MD00

Branch delay slot scheduling



Q. What to put in the delay slot?



Summary

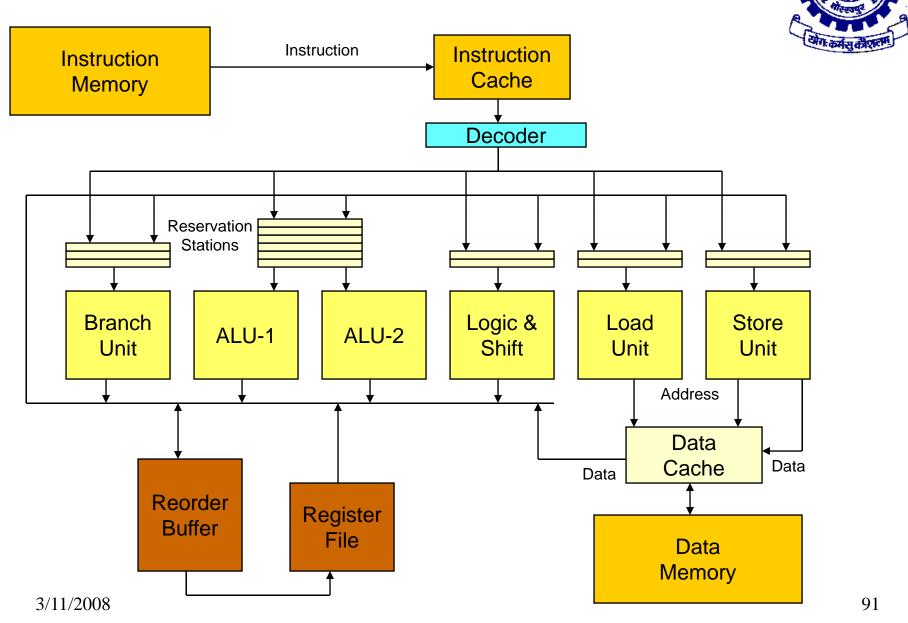


- Modern processors are (deeply) pipelined, to reduce T_{cycle} and aim at CPI = 1
- Hazards increase CPI
- Several software and hardware measure to avoid or reduce hazards are taken

Not discussed, but important developments:

- Multi-issue further reduces CPI
- Branch prediction to avoid high branch penalties
- Dynamic scheduling
- In all cases: a scheduling compiler needed

Superscalar Concept



Superscalar Execution



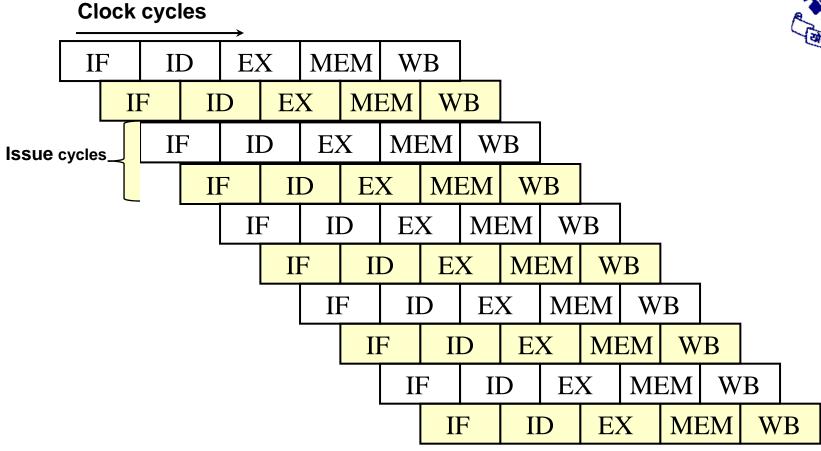
Clock cycles

Issue cycles_

		→						
IF	ID	EX	MEM	WB				
IF	ID	EX	MEM	WB				
	IF	ID	EX	MEM	WB			
	IF	ID	EX	MEM	WB			
		IF	ID	EX	MEM	WB		
		IF	ID	EX	MEM	WB		
			IF	ID	EX	MEM	WB	
			IF	ID	EX	MEM	WB	
				IF	ID	EX	MEM	WB
			'	IF	ID	EX	MEM	WB

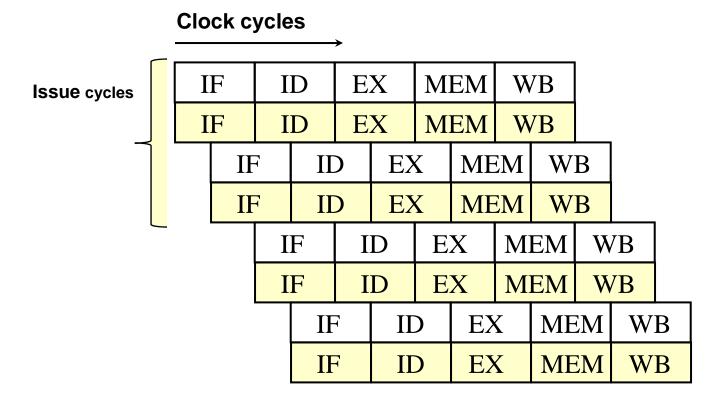
Superpipelined Execution





Superpipelined Superscalar Execution





Superscalar Issues



- How to fetch multiple instructions in time (across basic block boundaries)?
- Predicting branches
- Non-blocking memory system
- Tune #resources(FUs, ports, entries, etc.)
- Handling dependencies
- How to support precise interrupts?
- How to recover from mis-predicted branch path?
- For the latter two issues we need to look at sequential lookahead and architectural state

- Superscalar processor organization:
 - simple pipeline: IF, EX, WB
 - fetches 2 instructions each cycle
 - 2 Id/st units, dual-ported memory; 2 FP adders; 1 FP multiplier
 - Instruction window (buffer between IF and EX stage) is of size 2
 - FP ld/st takes 1 cc; FP +/- takes 2 cc; FP * takes 4 cc; FP / takes 8 cc

Cycle		1	2	3	4	5
L.D	F6,32(R2)					
L.D	F2,48(R3)					
MUL.D	F0,F2,F4					
SUB.D	F8,F2,F6					
DIV.D	F10,F0,F6					
ADD.D	F6,F8,F2					
MUL.D	F12,F2,F4					

3/11/2008

6



6

- Superscalar processor organization:
 - simple pipeline: IF, EX, WB
 - fetches 2 instructions each cycle
 - 2 ld/st units, dual-ported memory; 2 FP adders; 1 FP multiplier
 - Instruction window (buffer between IF and EX stage) is of size 2
 - FP ld/st takes 1 cc; FP +/- takes 2 cc; FP * takes 4 cc; FP / takes 8 cc

```
Cycle 1 2 3 4 5
L.D F6,32(R2) IF
L.D F2,48(R3) IF
MUL.D F0,F2,F4
SUB.D F8,F2,F6
DIV.D F10,F0,F6
ADD.D F6,F8,F2
```

3/11/2008

MUL.D F12,F2,F4



- Superscalar processor organization:
 - simple pipeline: IF, EX, WB
 - fetches 2 instructions each cycle
 - 2 Id/st units, dual-ported memory; 2 FP adders; 1 FP multiplier
 - Instruction window (buffer between IF and EX stage) is of size 2
 - FP ld/st takes 1 cc; FP +/- takes 2 cc; FP * takes 4 cc; FP / takes 8 cc

Cycle		1	2	3	4	5	6
L.D	F6,32(R2)	IF	EX				
L.D	F2,48(R3)	IF	EX				
MUL.D	F0,F2,F4		IF				
SUB.D	F8,F2,F6		IF				
DIV.D	F10,F0,F6						
ADD.D	F6,F8,F2						

3/11/2008

F12,F2,F4

MUL.D



- Superscalar processor organization:
 - simple pipeline: IF, EX, WB
 - fetches 2 instructions each cycle
 - 2 Id/st units, dual-ported memory; 2 FP adders; 1 FP multiplier
 - Instruction window (buffer between IF and EX stage) is of size 2
 - FP ld/st takes 1 cc; FP +/- takes 2 cc; FP * takes 4 cc; FP / takes 8 cc

Cycle		1	2	3	4	5	6	7
L.D	F6,32(R2)	IF	EX	WB				
L.D	F2,48(R3)	IF	EX	WB				
MUL.D	F0,F2,F4		IF	EX				
SUB.D	F8,F2,F6		IF	EX				
DIV.D	F10,F0,F6			IF				
ADD.D	F6,F8,F2			IF				
MUL.D	F12,F2,F4							



- Superscalar processor organization:
 - simple pipeline: IF, EX, WB
 - fetches 2 instructions each cycle
 - 2 Id/st units, dual-ported memory; 2 FP adders; 1 FP multiplier
 - Instruction window (buffer between IF and EX stage) is of size 2
 - FP ld/st takes 1 cc; FP +/- takes 2 cc; FP * takes 4 cc; FP / takes 8 cc

Cycle		1	2	3	4	5	6	7	
L.D	F6,32(R2)	IF	EX	WB					
L.D	F2,48(R3)	IF	EX	WB					
MUL.D	F0,F2,F4		IF	EX	EX				
SUB.D	F8,F2,F6		IF	EX	EX				
DIV.D	F10,F0,F6			IF	stall bed				
ADD.D	F6,F8,F2			IF	of data	dep.			
MUL.D	F12,F2,F4			canno	cannot be fetched because window fu				



- Superscalar processor organization:
 - simple pipeline: IF, EX, WB
 - fetches 2 instructions each cycle
 - 2 Id/st units, dual-ported memory; 2 FP adders; 1 FP multiplier
 - Instruction window (buffer between IF and EX stage) is of size 2
 - FP ld/st takes 1 cc; FP +/- takes 2 cc; FP * takes 4 cc; FP / takes 8 cc

Cycle		1	2	3	4	5	6	7
L.D	F6,32(R2)	IF	EX	WB				
L.D	F2,48(R3)	IF	EX	WB				
MUL.D	F0,F2,F4		IF	EX	EX	EX		
SUB.D	F8,F2,F6		IF	EX	EX	WB		
DIV.D	F10,F0,F6			IF				
ADD.D	F6,F8,F2			IF		EX		
MUL.D	F12,F2,F4					IF		



- Superscalar processor organization:
 - simple pipeline: IF, EX, WB
 - fetches 2 instructions each cycle
 - 2 Id/st units, dual-ported memory; 2 FP adders; 1 FP multiplier
 - Instruction window (buffer between IF and EX stage) is of size 2
 - FP Id/st takes 1 cc; FP +/- takes 2 cc; FP * takes 4 cc; FP / takes 8 cc

Cycle		1	2	3	4	5	6	7
L.D	F6,32(R2)	IF	EX	WB				
L.D	F2,48(R3)	IF	EX	WB				
MUL.D	F0,F2,F4		IF	EX	EX	EX	EX	
SUB.D	F8,F2,F6		IF	EX	EX	WB		
DIV.D	F10,F0,F6			IF				
ADD.D	F6,F8,F2			IF		EX	EX	
MUL.D	F12,F2,F4					IF	cannot exec	ute
3/11/20	08						structural haz	ard



- Superscalar processor organization:
 - simple pipeline: IF, EX, WB
 - fetches 2 instructions each cycle
 - 2 Id/st units, dual-ported memory; 2 FP adders; 1 FP multiplier
 - Instruction window (buffer between IF and EX stage) is of size 2
 - FP ld/st takes 1 cc; FP +/- takes 2 cc; FP * takes 4 cc; FP / takes 8 cc

Cycle		1	2	3	4	5	6	7
L.D	F6,32(R2)	IF	EX	WB				
L.D	F2,48(R3)	IF	EX	WB				
MUL.D	F0,F2,F4		IF	EX	EX	EX	EX	WB
SUB.D	F8,F2,F6		IF	EX	EX	WB		
DIV.D	F10,F0,F6			IF				EX
ADD.D	F6,F8,F2			IF		EX	EX	WB
MUL.D	F12,F2,F4					IF		?

Register Renaming



- A technique to eliminate anti- and output dependencies
- Can be implemented
 - by the compiler
 - advantage: low cost
 - disadvantage: "old" codes perform poorly
 - in hardware
 - advantage: binary compatibility
 - disadvantage: extra hardware needed
- We describe general idea

Register Renaming

- there's a *physical register file* larger than *logical register file*
- mapping table associates logical registers with physical register
- when an instruction is decoded
 - its physical source registers are obtained from mapping table
 - its physical destination register is obtained from a *free list*
 - · mapping table is updated

before: add	l r3,r3	, 4 after: add	R2,R1,	4
mapping table: r0	R8	mapping table: r0	R8	
r1	R7	r1	R7	
r2	R5	r2	R5	
r3	R1	r3	R2	
r4	R9	r4	R9	
free list:	R2 R6	free list:	R6	

Eliminating False Dependencies



- How register renaming eliminates false dependencies:
- Before:
 - addi r1, r2, 1
 - addi r2, r0, 0
 - addi r1, r0, 1
- After (free list: R7, R8, R9)
 - addi R7, R5, 1
 - addi R8, R0, 0
 - addi R9, R0, 1

Limitations of Multiple-Issue Processors



- Available ILP is limited (we're not programming with parallelism in mind)
- Hardware cost
 - adding more functional units is easy
 - more memory ports and register ports needed
 - dependency check needs O(n2) comparisons
- Limitations of VLIW processors
 - Loop unrolling increases code size
 - Unfilled slots waste bits
 - Cache miss stalls pipeline
 - Research topic: scheduling loads
 - Binary incompatibility (not EPIC)

Measuring available ILP: How?



- Using existing compiler
- Using trace analysis
 - Track all the real data dependencies (RaWs) of instructions from issue window
 - register dependence
 - memory dependence
 - Check for correct branch prediction
 - if prediction correct continue
 - if wrong, flush schedule and start in next cycle

Trace analysis

Program

Compiled code

How parallel can this code be executed?

Trace

set r1,0

set r2,3

set r3,&A

r1,0(r3)

add r1, r1, 1

add r3, r3, 4

brne r1,r2,Loop

st r1,0(r3)

add r1,r1,1

add r3, r3, 4

brne r1,r2,Loop

st r1,0(r3)

add r1,r1,1

add r3,r3,4

brne r1,r2,Loop

add r1, r5, 3

Trace analysis



Parallel Trace

```
set r1,0 set r2,3 set r3,&A
st r1,0(r3) add r1,r1,1 add r3,r3,4
st r1,0(r3) add r1,r1,1 add r3,r3,4 brne r1,r2,Loop
st r1,0(r3) add r1,r1,1 add r3,r3,4 brne r1,r2,Loop
brne r1,r2,Loop
add r1,r5,3
```

Speedup =
$$L_{parallel} / L_{serial} = 16 / 6 = 2.7$$

Is this the maximum?

Ideal Processor



Assumptions for ideal/perfect processor:

- 1. Register renaming infinite number of virtual registers => all register WAW & WAR hazards avoided
- 2. **Branch and Jump prediction** Perfect => all program instructions available for execution
- 3. *Memory-address alias analysis* addresses are known. A store can be moved before a load provided addresses not equal

Also:

- unlimited number of instructions issued/cycle (unlimited resources), and
- unlimited instruction window
- perfect caches
- 1 cycle latency for all instructions (FP *,/)

Programs were compiled using MIPS compiler with maximum optimization level